

ITCS-665 Natural Language Processing - Midterm Project

Sentiment Analysis and Toxicity Detection for Text Messages

Thant Zin Myo (6837929)

Contents

1. Introduction.....	2
2. Goal & Objectives.....	2
3. Learning Outcomes.....	3
4. Dataset Description.....	3
5. Exploratory Data Analysis	4
6. Preprocessing.....	5
6A. Cleaning	5
6B. Tokenization & Lemmatization.....	5
6C. Feature Engineering	5
7. Model Selection	5
8. Modeling.....	6
8.1 Full Process.....	6
9. Evaluation	7
10. Comparative Analysis	7
11. Ethical Considerations	8
12. Conclusion and Future Work	8
13. References.....	9
14. Acknowledgment	9

Table of Figures

Figure 1 Class Distribution	3
Figure 2 Data analysis with numbers.....	4
Figure 3 Word distribution by class	4
Figure 4 Word count percentiles and outlier	4
Figure 5 Text Pattern Analysis for Cleaning.....	4
Figure 6 Full Process	6
Figure 7 Comparison Summary	7
Figure 8 Confusion Matrices	7
Figure 9 Model Performance Comparison.....	8

1. Introduction

My current company is developing a live streaming application where viewers can interact in real time by sending chat messages and stickers, similar to popular platforms like Twitch or TikTok. These interactive features make streams more lively and enjoyable, as the audience can share thoughts, reactions, and support instantly.

However, when many people join a live session, the chat often includes toxic or hateful comments. Such comments can harm the streamer, create a negative environment, and even push other viewers away. To maintain a positive community, there is a need for a system that can automatically detect and filter harmful messages.

This project focuses on building a toxic comment classification model. The model will analyze each message and decide whether it is safe (non-toxic) or harmful (toxic). While the main use case is live streaming, such a system is also valuable for online games, forums, and social media where open chat is common.

The overall goal is to create a safer and more inclusive environment that protects users while still allowing them to freely express themselves with messages and stickers.

2. Goal & Objectives

- Apply knowledge of text preprocessing and classification techniques.
- Design and implement a text classification.
- Evaluate and compare models with and without Lexicons helps.

3. Learning Outcomes

By the end of this project, we can:

- Preprocess and prepare raw text data for classification.
- Apply and compare classification approaches.
- Analyze performance metrics and interpret model behavior.
- Communicate findings effectively through documentation and visualization.

4. Dataset Description

For this project, we are using the dataset from **Kaggle's Jigsaw Toxic Comment Classification Challenge**. This dataset has comments collected from Wikipedia pages. Each comment is labeled into six toxic categories: **toxic, severe toxic, obscene, threat, insult, and identity hate**. To make the problem easier we combined all of them into one label as we will do **Sentiment Analysis**:

- **toxic (1)** -> if the comment belongs to any of the six categories
- **non-toxic (0)** -> if the comment does not belong to any category

Dataset details:

- **Size:** 159,571 comments
- **Class balance:** about **90% non-toxic** and **10% toxic**
- **Format:** CSV file with comment text and multiple label columns




Figure 1 Class Distribution

In addition to the dataset, we will also use **Five lexicon hate word lists** ([luis von ahn-badwords.txt](#), [ldnoobw-bad-words](#), [pmathur5k10 -hinglish-offensive-texts](#), [orthrus-lexicontoxic](#), [multilingual toxic lexicon](#)). These lexicons are collections of words that are often used in toxic or hateful comments. By combining the dataset with the lexicon approach, we can compare how well machine learning models work alone, and with the help of lexicons.

5. Exploratory Data Analysis

- Label Distribution:** As we can see in Fig-1 Distribution is quite imbalanced around 90- (Non-toxic: 143k, Toxic: 16k). and average length with 394 characters and around 67 words per comments. We also detect some text with HTML, Quoted Text, Number and URL which are necessary for cleaning.

```

EXPLORATORY DATA ANALYSIS
Label Distribution:
Non-toxic (0): 143,346 (89.8%)
Toxic (1): 16,225 (10.2%)
Average text length: 394 characters
Average word count: 67 words
  
```

Figure 2 Data analysis with numbers




Figure 3 Word distribution by class




Figure 4 Word count percentiles and outlier

- Patterns detected:** 5,151 comments with URLs, 49,996 with quotes, 51,212 with numbers.




Figure 5 Text Pattern Analysis for Cleaning

6. Preprocessing

6A. Cleaning

The cleaning pipeline included:

- Lowercasing text
- Removing numbers, quotes, URLs, and HTML tags
- Expanding contractions (e.g., *don't* -> *do not*) with the help of [andrewbury contractions sources](#)
- Removing punctuation (keeping only alphabetic tokens)
- Stripping extra whitespace

6B. Tokenization & Lemmatization

- **Tokenization:** Splitting text into words using NLTK.
- **Stop word removal:** Common English words removed (e.g., “the”, “and”).
- **Lemmatization:** Words reduced to their base form (e.g., *running* -> *run*).

6C. Feature Engineering

We experimented with:

- **TF-IDF vectorization** with unigrams and bigrams (max features=10,000).
- **Count Vectorizer** with unigrams and bigrams (max features=10,000) + (4760 words but remove duplicate base on feature out-put of the model).

7. Model Selection

We focused on simple but effective classifiers called **Multinomial Naive Bayes (MNB)** and **Logistic Regression (LR)**. It is often used in text classification tasks.

We tested 4 different models with different setting:

- **Model 1: Multinomial Naive Bayes**
 - **Setting:** TF-IDF (feature =10000), Cleaning, Tokenization, Lemmatization, Remove stop word
- **Model 2: Multinomial Naive Bayes**
 - **Setting:** Bag of words (feature =10000+ lexicon), Cleaning, Tokenization & Lemmatization, Remove stop word

- **Model 3: Logistic Regression**
 - **Setting: TF-IDF (feature =10000), Cleaning, Tokenization & Lemmatization, Remove stop word**
- **Model 4: Logistic Regression**
 - **Setting: Lexicon, Bag of words (feature =10000+ lexicon), Cleaning, Tokenization & Lemmatization, Remove stop word**

8. Modeling

Train/Test Split

We divided the dataset into two parts:

For without lexicons Dataset we split into

- 80% for training (127,656 samples)
- 20% for testing (31,915 samples)

For with lexicons Dataset we split into

- 80% for training (127,656 samples)+ (lexicon word as text and label as 1) = 132, 416 sample
- 20% for testing (31,915 samples)

8.1 Full Process




Figure 6 Full Process

9. Evaluation

We measured accuracy, precision, recall, F1-score, and confusion matrix values. Naive Bayes with TF-IDF reached 94.32% accuracy, 93.82% precision, 47.27% recall, and an F1-score of 62.87%. Adding lexicons and Changing TF-IDF to Bag of words improved recall to 54.33% and F1 to 67.17%. Logistic Regression with TF-IDF gave 93.22% accuracy, 64.25% precision, 75.10% recall, and F1 of 69.25%. With lexicons and changing to TF-IDF Bag of words improved slightly, reaching 93.75% accuracy, 67.66% precision, 73.84% recall, and F1 of 70.62%.

Overall, Naive Bayes was precise but missed many toxic comments, while Logistic Regression detected more toxic cases and offered a better balance. Lexicons + bag of words models gave small improvements in both models.

FINAL COMPARISON SUMMARY

Model	Accuracy	F1-Score	Precision	Recall	True Positives	True Negatives	False Positives	False Negatives
1: MNB+TFIDF	0.9432	0.6287	0.9382	0.4727	1534	28569	101	1711
2: MNB+BOW+LEX	0.946	0.6717	0.8797	0.5433	1763	28429	241	1482
3: LR+TFIDF	0.9322	0.6925	0.6425	0.751	2437	27314	1356	808
4: MNB+BOW+LEX	0.9375	0.7062	0.6766	0.7384	2396	27525	1145	849

Figure 7 Comparison Summary

Individual Confusion Matrices							
1: MNB+TFIDF		2: MNB+BOW+LEX		3: LR+TFIDF		4: MNB+BOW+LEX	
TP	1534	TP	1763	TP	2437	TP	2396
	28569		28429		27314		27525
TN		TN		TN		TN	
FP	101	FP	241	FP	1356	FP	1145
	1711		1482		808		849
FN		FN		FN		FN	

Figure 8 Confusion Matrices

10. Comparative Analysis

Naive Bayes worked best for precision, meaning it rarely misclassified safe comments as toxic, but it missed many toxic ones. Logistic Regression was more balanced and achieved higher recall, which is more useful for real-time moderation. Adding lexicons (about **4,760 hate words**) helped both models by catching more toxic words, though at the cost of some extra false positives. The best overall model was Logistic Regression with bag of words and lexicons.




Figure 9 Model Performance Comparison

11. Ethical Considerations

Toxic comment detection raises fairness and bias concerns. Since the dataset and lexicons are mainly in English, the models may fail on other languages or unfairly flag dialects and slang. Overly strict filtering can also block harmless jokes or expressions. Because of this, such systems should support multiple languages and always include report or reclaim function so that we can moderate to avoid unfair censorship.

12. Conclusion and Future Work

For future work, more advanced models or deep learning models like BERT should be tested, cross-validation should be used for stronger evaluation, and error analysis should be done to see which comments are missed. Expanding to multiple languages and combining with human moderators will make the system fairer and more reliable.

Currently, the best model achieved an F1-score of around **0.7**, which shows good progress but also highlights that further improvement is still needed. Additional training with larger and more diverse datasets, along with fine-tuning on real-world chat data, will help improve overall performance.

13. References

- Kaggle. (2018). *Jigsaw Toxic Comment Classification Challenge*.
<https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge>
- von Ahn, L. (n.d.). *CMU Offensive/Profanity Word List*. Carnegie Mellon University.
<https://www.cs.cmu.edu/~biglou/resources/>
- LDNOOBW. (n.d.). *List of Dirty, Naughty, Obscene, and Otherwise Bad Words*. GitHub.
<https://github.com/LDNOOBW/List-of-Dirty-Naughty-Obscene-and-Otherwise-Bad-Words>
- Mathur, P. (n.d.). *Hinglish Offensive Text Classification – Profanity Lexicon*. GitHub.
<https://github.com/pmathur5k10/Hinglish-Offensive-Text-Classification>
- Orthrus Lexicon Project. (n.d.). *Toxic Words Dictionary*. GitHub.
<https://github.com/Orthrus-Lexicon/Toxic>
- Hugging Face. (n.d.). *Multilingual Toxic Lexicon – TextDetox Dataset*.
https://huggingface.co/datasets/textdetox/multilingual_toxic_lexicon
- Bury, A. (n.d.). *English Contractions List for NLP Tasks*. GitHub.
<https://github.com/andrewbury/contractions/blob/master/contractions.txt>

14. Acknowledgment

All modeling and analysis were carried out by the author, while ChatGPT (current version GPT-5) was only used to improve grammar and readability.

The Jupyter Notebook and training dataset can be downloaded from GitHub:

<https://github.com/karengsengli/NLP-MidTermProject>